

Chapter 6: Data Cleaning, Loops, apply functions

M Affouf

1/8/2018

part1

Data

- ▶ We will be using multiple data sets in this lecture:
 - ▶ Salary, Monument, Circulator, and Restaurant from OpenBaltimore: <https://data.baltimorecity.gov/browse?limitTo=datasets>
 - ▶ Gap Minder - very interesting way of viewing longitudinal data
 - ▶ Data is here - <http://www.gapminder.org/data/>
 - ▶ http://spreadsheets.google.com/pub?key=rMsQHawT0bBb6_U2ESjKXYw&output=xls

Data Cleaning

In general, data cleaning is a process of investigating your data for inaccuracies, or recoding it in a way that makes it more manageable.

MOST IMPORTANT RULE - LOOK AT YOUR DATA!

Again - `table`, `summarize`, `is.na`, `any`, `all` are useful.

Dealing with Missing Data

Missing data types

One of the most important aspects of data cleaning is missing values.

Types of “missing” data:

- ▶ NA - general missing data
- ▶ NaN - stands for “**N**ot **a** **N**umber”, happens when you do $0/0$.
- ▶ Inf and -Inf - Infinity, happens when you take a positive number (or negative number) by 0.

Finding Missing data

Each missing data type has a function that returns TRUE if the data is missing:

- ▶ NA - `is.na`
- ▶ NaN - `is.nan`
- ▶ Inf and -Inf - `is.infinite`
- ▶ `is.finite` returns FALSE for all missing data and TRUE for non-missing
- ▶ `complete.cases` on a `data.frame/matrix` returns TRUE if all values in that row of the object are not missing.

Missing Data with Logicals

One important aspect (esp with subsetting) is that logical operations return NA for NA values. Think about it, the data could be > 2 or not we don't know, so R says there is no TRUE or FALSE, so that is missing:

```
x = c(0, NA, 2, 3, 4)
x > 2
```

```
## [1] FALSE      NA FALSE  TRUE  TRUE
```


Missing Data with Logicals

What to do? What if we want if $x > 2$ and x isn't NA?

Don't do $x \neq \text{NA}$, do $x > 2$ and x is NOT NA:

```
x != NA
```

```
## [1] NA NA NA NA NA
```

```
x > 2 & !is.na(x)
```

```
## [1] FALSE FALSE FALSE  TRUE  TRUE
```

Missing Data with Logicals

What about seeing if a value is equal to multiple values? You can do `(x == 1 | x == 2) & !is.na(x)`, but that is not efficient. Introduce the `%in%` operator:

```
(x == 0 | x == 2) # has NA
```

```
## [1] TRUE NA TRUE FALSE FALSE
```

```
(x == 0 | x == 2) & !is.na(x) # No NA
```

```
## [1] TRUE FALSE TRUE FALSE FALSE
```

```
x %in% c(0, 2) # NEVER has NA and returns logical
```

```
## [1] TRUE FALSE TRUE FALSE FALSE
```

Missing Data with Operations

Similarly with logicals, operations/arithmetic with NA will result in NAs:

```
x + 2
```

```
## [1] 2 NA 4 5 6
```

```
x * 2
```

```
## [1] 0 NA 4 6 8
```

Tables and Tabulations

Creating One-way Tables

Here we will use `table` to make tabulations of the data. Look at `?table` to see options for missing data.

```
table(x)
```

```
## x  
## 0 2 3 4  
## 1 1 1 1
```

```
table(x, useNA = "ifany")
```

```
## x  
##    0    2    3    4 <NA>  
##    1    1    1    1    1
```

Creating One-way Tables

You can set `useNA = "always"` to have it always have a column for NA

```
table(c(0, 1, 2, 3, 2, 3, 3, 2,2, 3),  
      useNA = "always")
```

```
##  
##      0      1      2      3 <NA>  
##      1      1      4      4      0
```

Creating Two-way Tables

A two-way table. If you pass in 2 vectors, `table` creates a 2-dimensional table.

```
tab <- table(c(0, 1, 2, 3, 2, 3, 3, 2, 2, 3),  
             c(0, 1, 2, 3, 2, 3, 3, 4, 4, 3),  
             useNA = "always")
```

Finding Row or Column Totals

`margin.table` finds the marginal sums of the table. `margin` is 1 for rows, 2 for columns in general in R. Here is the column sums of the table:

```
margin.table(tab, 2)
```

```
##  
##      0      1      2      3      4 <NA>  
##      1      1      2      4      2      0
```


Proportion Tables

`prop.table` finds the marginal proportions of the table. Think of it dividing the table by it's respective marginal totals. If `margin` not set, divides by overall total.

```
prop.table(tab)
```

```
##
##           0    1    2    3    4 <NA>
##  0      0.1 0.0 0.0 0.0 0.0  0.0
##  1      0.0 0.1 0.0 0.0 0.0  0.0
##  2      0.0 0.0 0.2 0.0 0.2  0.0
##  3      0.0 0.0 0.0 0.4 0.0  0.0
## <NA> 0.0 0.0 0.0 0.0 0.0  0.0
```

```
prop.table(tab,1)
```

```
##
##           0    1    2    3    4 <NA>
```

Download Salary FY2014 Data

From <https://data.baltimorecity.gov/City-Government/Baltimore-City-Employee-Salaries-FY2014/2j28-xzd7>

Read the CSV into R Sal:

```
Sal = read.csv("Baltimore_City_Employee_Salaries_FY2014.csv",  
               as.is = TRUE)
```

Checking for logical conditions

- ▶ `any()` - checks if there are any TRUEs
- ▶ `all()` - checks if ALL are true

```
head(Sal,2)
```

```
##              Name              JobTitle Agency
## 1  Aaron,Keontae E          AIDE BLUE CHIP  W0220
## 2 Aaron,Patricia G Facilities/Office Services II  A0303
##              Agency  HireDate AnnualSalary  GrossPay
## 1      Youth Summer  06/10/2013    $11310.00    $873.63
## 2 OED-Employment Dev  10/24/1979    $53428.00 $52868.38
```

```
any(is.na(Sal$Name)) # are there any NAs?
```

```
## [1] FALSE
```

Recoding Variables

Example of Recoding: base R

For example, let's say gender was coded as Male, M, m, Female, F, f. Using Excel to find all of these would be a matter of filtering and changing all by hand or using if statements.

In R, you can simply do something like:

```
# data$gender[data$gender %in%  
# c("Male", "M", "m")] <- "Male"
```

Example of Recoding with recode: car package

You can also recode a vector:

```
library(car, quietly = TRUE)
x = rep(c("Male", "M", "m", "f", "Female", "female" ),
        each = 3)
car::recode(x, "c('m', 'M', 'male') = 'Male';
              c('f', 'F', 'female') = 'Female';")
```

```
## [1] "Male" "Male" "Male" "Male" "Male" "Male"
## [8] "Male" "Male" "Female" "Female" "Female" "Female"
## [15] "Female" "Female" "Female" "Female"
```

Example of Recoding with revalue: plyr

You can also revalue a vector with the revalue command

```
library(plyr)
plyr::revalue(x, c("M" = "Male", "m" = "Male",
                  "f" = "Female", "female" = "Female"))

## [1] "Male" "Male" "Male" "Male" "Male" "Male"
## [8] "Male" "Male" "Female" "Female" "Female" "Female"
## [15] "Female" "Female" "Female" "Female"
```

Example of Cleaning: more complicated

Sometimes though, it's not so simple. That's where functions that find patterns come in very useful.

```
table(gender)
```

```
## gender
##      F FeMAle FEMALE      Fm      M      Ma      mAle      Male
##      75      82      74      89      89      79      87      89
##      Man      Woman
##      73      80
```


String functions

Pasting strings with paste and paste0

Paste can be very useful for joining vectors together:

```
paste("Visit", 1:5, sep = "_")
```

```
## [1] "Visit_1" "Visit_2" "Visit_3" "Visit_4" "Visit_5"
```

```
paste("Visit", 1:5, sep = "_", collapse = " ")
```

```
## [1] "Visit_1 Visit_2 Visit_3 Visit_4 Visit_5"
```

```
paste("To", "is going be the ", "we go to the store!", sep
```

```
## [1] "Today is going be the day we go to the store!"
```

and paste0 can be even simpler see ?paste0

```
paste0("Visit",1:5)
```

```
## [1] "Visit1" "Visit2" "Visit3" "Visit4" "Visit5"
```

Paste Depicting How Collapse Works

```
paste(1:5)
```

```
## [1] "1" "2" "3" "4" "5"
```

```
paste(1:5, collapse = " ")
```

```
## [1] "1 2 3 4 5"
```

Useful String Functions

Useful String functions

- ▶ `toupper()`, `tolower()` - uppercase or lowercase your data:
- ▶ `str_trim()` (in the `stringr` package) or `trimws` in base
 - ▶ will trim whitespace
- ▶ `nchar` - get the number of characters in a string
- ▶ `paste()` - paste strings together with a space
- ▶ `paste0` - paste strings together with no space as default

The stringr package

Like dplyr, the stringr package:

- ▶ Makes some things more intuitive
- ▶ Is different than base R
- ▶ Is used on forums for answers
- ▶ Has a standard format for most functions
 - ▶ the first argument is a string like first argument is a `data.frame` in dplyr

Splitting/Find/Replace and Regular Expressions

- ▶ R can do much more than find exact matches for a whole string
- ▶ Like Perl and other languages, it can use regular expressions.
- ▶ What are regular expressions?
 - ▶ Ways to search for specific strings
 - ▶ Can be very complicated or simple
 - ▶ Highly Useful - think “Find” on steroids

A bit on Regular Expressions

- ▶ `http:`
`//www.regular-expressions.info/reference.html`
- ▶ They can use to match a large number of strings in one statement
- ▶ `.` matches any single character
- ▶ `*` means repeat as many (even if 0) more times the last character
- ▶ `?` makes the last thing optional
- ▶ `^` matches start of vector `^a` - starts with "a"
- ▶ `$` matches end of vector `b$` - ends with "b"

Splitting Strings

Substringing

Very similar:

Base R

- ▶ `substr(x, start, stop)` - substrings from position start to position stop
- ▶ `strsplit(x, split)` - splits strings up - returns list!

`stringr`

- ▶ `str_sub(x, start, end)` - substrings from position start to position end
- ▶ `str_split(string, pattern)` - splits strings up - returns list!

Splitting String: base R

In base R, `strsplit` splits a vector on a string into a list

```
x <- c("I really", "like writing", "R code programs")  
y <- strsplit(x, split = " ") # returns a list  
y
```

```
## [[1]]  
## [1] "I"      "really"  
##  
## [[2]]  
## [1] "like"    "writing"  
##  
## [[3]]  
## [1] "R"      "code"    "programs"
```

Splitting String: stringr

stringr::str_split do the same thing:

```
library(stringr)
y2 <- str_split(x, " ") # returns a list
y2
```

```
## [[1]]
## [1] "I"      "really"
##
## [[2]]
## [1] "like"    "writing"
##
## [[3]]
## [1] "R"      "code"    "programs"
```

Using a fixed expression

One example case is when you want to split on a period “.”. In regular expressions `.` means **ANY** character, so

```
str_split("I.like.strings", ".")
```

```
## [[1]]
```

```
## [1] "" "" "" "" "" "" "" "" "" "" "" "" "" "" ""
```

```
str_split("I.like.strings", fixed("."))
```

```
## [[1]]
```

```
## [1] "I"      "like"   "strings"
```

Let's extract from y

```
suppressPackageStartupMessages(library(dplyr)) # must be loaded  
y[[2]]
```

```
## [1] "like"      "writing"
```

```
sapply(y, dplyr::first) # on the fly
```

```
## [1] "I"        "like" "R"
```

```
sapply(y, nth, 2) # on the fly
```

```
## [1] "really"    "writing" "code"
```

```
sapply(y, last) # on the fly
```

```
## [1] "really"    "writing"  "programs"
```

'Find' functions: base R

`grep`: `grep`, `grep1`, `regexpr` and `gregexpr` search for matches to argument `pattern` within each element of a character vector: they differ in the format of and amount of detail in the results.

`grep(pattern, x, fixed=FALSE)`, where:

- ▶ `pattern` = character string containing a regular expression to be matched in the given character vector.
- ▶ `x` = a character vector where matches are sought, or an object which can be coerced by `as.character` to a character vector.
- ▶ If `fixed=TRUE`, it will do exact matching for the phrase anywhere in the vector (regular find)

'Find' functions: `stringr`

`str_detect`, `str_subset`, `str_replace`, and `str_replace_all` search for matches to argument pattern within each element of a character vector: they differ in the format of and amount of detail in the results.

- ▶ `str_detect` - returns TRUE if pattern is found
- ▶ `str_subset` - returns only the strings which pattern were detected
 - ▶ convenient wrapper around `x[str_detect(x, pattern)]`
- ▶ `str_extract` - returns only strings which pattern were detected, but ONLY the pattern
- ▶ `str_replace` - replaces pattern with replacement the first time
- ▶ `str_replace_all` - replaces pattern with replacement as many times matched

'Find' functions: stringr compared to base R

Base R does not use these functions. Here is a “translator” of the `stringr` function to base R functions

- ▶ `str_detect` - similar to `grepl` (return logical)
- ▶ `grep(value = FALSE)` is similar to `which(str_detect())`
- ▶ `str_subset` - similar to `grep(value = TRUE)` - return value of matched
- ▶ `str_replace` - similar to `sub` - replace one time
- ▶ `str_replace_all` - similar to `gsub` - replace many times

Let's look at modifier for stringr

?modifiers

- ▶ fixed - match everything exactly
- ▶ regexp - default - uses **regular expressions**
- ▶ ignore_case is an option to not have to use tolower

Important Comparisons

Base R:

- ▶ Argument order is (pattern, x)
- ▶ Uses option (fixed = TRUE)

stringr

- ▶ Argument order is (string, pattern) aka (x, pattern)
- ▶ Uses function fixed(pattern)

'Find' functions: Finding Indices

These are the indices where the pattern match occurs:

```
grep("Rawlings", Sal$Name)
```

```
## [1] 13832 13833 13834 13835
```

```
which(grepl("Rawlings", Sal$Name))
```

```
## [1] 13832 13833 13834 13835
```

```
which(str_detect(Sal$Name, "Rawlings"))
```

```
## [1] 13832 13833 13834 13835
```

'Find' functions: Finding Logicals

These are the indices where the pattern match occurs:

```
head(grepl("Rawlings",Sal$Name))
```

```
## [1] FALSE FALSE FALSE FALSE FALSE FALSE
```

```
head(str_detect(Sal$Name, "Rawlings"))
```

```
## [1] FALSE FALSE FALSE FALSE FALSE FALSE
```

'Find' functions: finding values, base R

```
grep("Rawlings",Sal$Name,value=TRUE)
```

```
## [1] "Rawlings,Kellye A"          "Rawlings,MarqWell D"  
## [3] "Rawlings,Paula M"          "Rawlings-Blake,Stephanie C"
```

```
Sal[grep("Rawlings",Sal$Name),]
```

```
##                               Name                JobTitle Age  
## 13832      Rawlings,Kellye A EMERGENCY DISPATCHER  
## 13833      Rawlings,MarqWell D      AIDE BLUE CHIP  
## 13834      Rawlings,Paula M      COMMUNITY AIDE  
## 13835 Rawlings-Blake,Stephanie C      MAYOR  
##                               Agency  HireDate  AnnualSalary  Gro  
## 13832 M-R Info Technology  01/06/2003    $47980.00  $684  
## 13833      Youth Summer  06/15/2012    $11310.00    $5  
## 13834      R&P-Recreation  12/10/2007    $19802.00    $81  
## 13835      Mayors Office  12/07/1995   $163365.00 $1612
```

'Find' functions: finding values, stringr and dplyr

```
str_subset(Sal$Name, "Rawlings")
```

```
## [1] "Rawlings,Kellye A"          "Rawlings,MarqWell D"  
## [3] "Rawlings,Paula M"          "Rawlings-Blake,Stephanie C"
```

```
Sal %>% filter(str_detect(Name, "Rawlings"))
```

```
##           Name           JobTitle Agency  
## 1 Rawlings,Kellye A EMERGENCY DISPATCHER  A403  
## 2 Rawlings,MarqWell D AIDE BLUE CHIP W023  
## 3 Rawlings,Paula M COMMUNITY AIDE A040  
## 4 Rawlings-Blake,Stephanie C MAYOR A010  
##           Agency HireDate AnnualSalary GrossPay  
## 1 M-R Info Technology 01/06/2003 $47980.00 $68426.7  
## 2 Youth Summer 06/15/2012 $11310.00 $507.5  
## 3 R&P-Recreation 12/10/2007 $19802.00 $8195.7  
## 4 Mayors Office 12/07/1995 $163365.00 $161219.2
```

Showing difference in str_extract

str_extract extracts just the matched string

```
ss = str_extract(Sal$Name, "Rawling")  
head(ss)
```

```
## [1] NA NA NA NA NA NA
```

```
ss[ !is.na(ss)]
```

```
## [1] "Rawling" "Rawling" "Rawling" "Rawling"
```

Showing difference in str_extract and str_extract_all

str_extract_all extracts all the matched strings

```
head(str_extract(Sal$AgencyID, "\\d"))
```

```
## [1] "0" "0" "2" "6" "9" "4"
```

```
head(str_extract_all(Sal$AgencyID, "\\d"), 2)
```

```
## [[1]]
```

```
## [1] "0" "2" "2" "0" "0"
```

```
##
```

```
## [[2]]
```

```
## [1] "0" "3" "0" "3" "1"
```


Using Regular Expressions

- ▶ Look for any name that starts with:
 - ▶ Payne at the beginning,
 - ▶ Leonard and then an S
 - ▶ Spence then capital C

```
head(grep("^Payne.*", x = Sal$Name, value = TRUE), 3)
```

```
## [1] "Payne El,Jackie"          "Payne Johnson,Nickole A"  
## [3] "Payne,Chanel"
```

```
head(grep("Leonard.?S", x = Sal$Name, value = TRUE))
```

```
## [1] "Payne,Leonard S"         "Szumlanski,Leonard S"
```

```
head(grep("Spence.*C.*", x = Sal$Name, value = TRUE))
```

```
## [1] "Greene,Spencer C"        "Spencer,Charles A"    "Spencer  
## [4] "Spencer,Clarence W"     "Spencer,Michael C"
```

Using Regular Expressions: stringr

```
head(str_subset( Sal$Name, "^Payne.*"), 3)
```

```
## [1] "Payne El,Jackie"          "Payne Johnson,Nickole A"  
## [3] "Payne,Chanel"
```

```
head(str_subset( Sal$Name, "Leonard.?S"))
```

```
## [1] "Payne,Leonard S"        "Szumlanski,Leonard S"
```

```
head(str_subset( Sal$Name, "Spence.*C.*"))
```

```
## [1] "Greene,Spencer C"      "Spencer,Charles A"    "Spencer  
## [4] "Spencer,Clarence W"    "Spencer,Michael C"
```

Replace

Let's say we wanted to sort the data set by Annual Salary:

```
class(Sal$AnnualSalary)
```

```
## [1] "character"
```

```
sort(c("1", "2", "10")) # not sort correctly (order simply)
```

```
## [1] "1" "10" "2"
```

```
order(c("1", "2", "10"))
```

```
## [1] 1 3 2
```

Replace

So we must change the annual pay into a numeric:

```
head(Sal$AnnualSalary, 4)
```

```
## [1] "$11310.00" "$53428.00" "$68300.00" "$62000.00"
```

```
head(as.numeric(Sal$AnnualSalary), 4)
```

```
## Warning in head(as.numeric(Sal$AnnualSalary), 4): NAs in  
## coercion
```

```
## [1] NA NA NA NA
```

R didn't like the \$ so it thought turned them all to NA.

sub() and gsub() can do the replacing part in base R.

Replacing and subbing

Now we can replace the \$ with nothing (used fixed=TRUE because \$ means ending):

```
Sal$AnnualSalary <- as.numeric(gsub(pattern = "$", replacement = "",  
                                   Sal$AnnualSalary, fixed=TRUE))  
Sal <- Sal[order(Sal$AnnualSalary, decreasing=TRUE), ]  
Sal[1:5, c("Name", "AnnualSalary", "JobTitle")]
```

##		Name	AnnualSalary	JobTitle
## 1222	Bernstein, Gregg L	238772	STATE'S ATTORNEY	
## 3175	Charles, Ronnie E	200000	EXECUTIVE LEVEL 1	
## 985	Batts, Anthony W	193800	EXECUTIVE LEVEL 1	
## 1343	Black, Harry E	190000	EXECUTIVE LEVEL 1	
## 16352	Swift, Michael	187200	CONTRACT SERV SPEC	

Replacing and subbing: stringr

We can do the same thing (with 2 piping operations!) in dplyr

```
dplyr_sal = Sal
dplyr_sal = dplyr_sal %>% mutate(
  AnnualSalary = AnnualSalary %>%
    str_replace(
      fixed("$"),
      "") %>%
    as.numeric() %>%
    arrange(desc(AnnualSalary))
check_Sal = Sal
rownames(check_Sal) = NULL
all.equal(check_Sal, dplyr_sal)
```

```
## [1] TRUE
```

Part 2

Agenda

- ▶ A common data cleaning task
- ▶ If-else statements
- ▶ For/while loops to iterate over data
- ▶ `apply()`, `lapply()`, `sapply()`, `tapply()`
- ▶ `with()` to specify scope

A common problem

- ▶ One of the most common problems you'll encounter when importing manually-entered data is inconsistent data types within columns
- ▶ For a simple example, let's look at TVhours column in a messy version of the survey data from Lecture 2

```
survey.messy <- read.csv("survey_messy.csv", header=TRUE)
survey.messy$TVhours
```

```
## [1] 4 ~5 2 5or so
## [6] 15 3 2ish 0
## [11] none 7 3 6.5
## [16] 0 gjkhgs 3 0
## [21] 3 4 10 2.5 (a mo
## [26] 6 zero 0 2
## [31] 2 4 4 0
## [36] 0 2
## 18 Levels: ~5 0 10 15 2 2.5 (a movie) 2ish 3 4 5 (netfl
```

What's happening?

```
str(survey.messy)
```

```
## 'data.frame':    37 obs. of  6 variables:
## $ Program      : Factor w/ 3 levels "MISM","Other",...
## $ PriorExp     : Factor w/ 3 levels "Extensive experi
## $ Rexperience   : Factor w/ 3 levels "Basic competence
## $ OperatingSystem: Factor w/ 2 levels "Mac OS X","Wind
## $ TVhours      : Factor w/ 18 levels "~5","0","10",...
## $ Editor       : Factor w/ 5 levels "Google Docs",...
```

- ▶ Several of the entries have non-numeric values in them (they contain strings)
- ▶ As a result, TVhours is being imported as factor

Attempt at a fix

- What if we just try to cast it back to numeric?

```
tv.hours.messy <- survey.messy$TVhours  
tv.hours.messy
```

```
## [1] 4 ~5 2 5or so  
## [6] 15 3 2ish 0  
## [11] none 7 3 6.5  
## [16] 0 gjkhgs 3 0  
## [21] 3 4 10 2.5 (a movie)  
## [26] 6 zero 0 2  
## [31] 2 4 4 0  
## [36] 0 2  
## 18 Levels: ~5 0 10 15 2 2.5 (a movie) 2ish 3 4 5 (netfl
```

```
as.numeric(tv.hours.messy)
```

```
## [1] 9 1 5 11 8 4 8 7 2 5 17 14 8 13 2 2 16
```

That didn't work...

```
tv.hours.messy  
as.numeric(tv.hours.messy)
```

```
## [1] 4 ~5 2 5or so  
## [6] 15 3 2ish 0  
## [11] none 7 3 6.5  
## [16] 0 gjkhgs 3 0  
## [21] 3 4 10 2.5 (a mo  
## [26] 6 zero 0 2  
## [31] 2 4 4 0  
## [36] 0 2  
## 18 Levels: ~5 0 10 15 2 2.5 (a movie) 2ish 3 4 5 (netfl  
  
## [1] 9 1 5 11 8 4 8 7 2 5 17 14 8 13 2 2 16  
## [24] 6 15 12 18 2 5 9 5 9 9 2 10 2 5
```

- This just converted all the values into the integer-coded levels of the factor

Something that does work

- Consider the following simple example

```
num.vec <- c(3.1, 2.5)  
as.factor(num.vec)
```

```
## [1] 3.1 2.5  
## Levels: 2.5 3.1
```

```
as.numeric(as.factor(num.vec))
```

```
## [1] 2 1
```

```
as.numeric(as.character(as.factor(num.vec)))
```

```
## [1] 3.1 2.5
```

If we take a number that's being coded as a factor and first turn it into a character string, then converting the string to a numeric gets back the number

Back to the corrupted TVhours column

```
as.character(tv.hours.messy)
```

```
## [1] "4"           "~5"          "2"           "50"
## [5] "3"           "15"         "3"           "23"
## [9] "0"           "2"          "none"        "7"
## [13] "3"           "6.5"        "0"           "0"
## [17] "gjkhgs"      "3"          "0"           "4"
## [21] "3"           "4"          "10"          "2"
## [25] "8"           "6"          "zero"        "0"
## [29] "2"           "4"          "2"           "4"
## [33] "4"           "0"          "5 (netflix)" "0"
## [37] "2"
```

```
as.numeric(as.character(tv.hours.messy))
```

```
## Warning: NAs introduced by coercion
```

```
## [1] 4.0 NA 2.0 NA 3.0 15.0 3.0 NA 0.0 2.0
```

A small improvement

- ▶ All the corrupted cells now appear as NA, which is R's missing indicator
- ▶ We can do a little better by cleaning up the vector once we get it to character form

```
tv.hours.strings <- as.character(tv.hours.messy)
tv.hours.strings
```

```
## [1] "4"           "~5"          "2"           "50"
## [5] "3"           "15"         "3"           "21"
## [9] "0"           "2"          "none"        "7"
## [13] "3"           "6.5"        "0"           "0"
## [17] "gjkhgs"      "3"          "0"           "4"
## [21] "3"           "4"          "10"          "2"
## [25] "8"           "6"          "zero"        "0"
## [29] "2"           "4"          "2"           "4"
## [33] "4"           "0"          "5 (netflix)" "0"
## [37] "2"
```

Deleting non-numeric (or .) characters

```
tv.hours.strings
```

```
## [1] "4"           "~5"          "2"           "50"
## [5] "3"           "15"         "3"           "23"
## [9] "0"           "2"          "none"        "7"
## [13] "3"           "6.5"        "0"           "0"
## [17] "gjkhgs"      "3"          "0"           "4"
## [21] "3"           "4"          "10"          "2"
## [25] "8"           "6"          "zero"        "0"
## [29] "2"           "4"          "2"           "4"
## [33] "4"           "0"          "5 (netflix)" "0"
## [37] "2"
```

```
# Use gsub() to replace everything except digits and '.' with  
gsub("[^0-9.]", "", tv.hours.strings)
```

```
## [1] "4" "5" "2" "5" "3" "15" "3" "2" "0"
## [12] "7" "3" "6.5" "0" "0" "" "3" "0" "4"
```


The final product

```
tv.hours.messy[1:30]
```

```
## [1] 4 ~5 2 5or so
## [6] 15 3 2ish 0
## [11] none 7 3 6.5
## [16] 0 gjkhgs 3 0
## [21] 3 4 10 2.5 (a movie)
## [26] 6 zero 0 2
## 18 Levels: ~5 0 10 15 2 2.5 (a movie) 2ish 3 4 5 (netfl)
```

```
tv.hours.clean <- as.numeric(gsub("[^0-9.]", "", tv.hours.messy))
tv.hours.clean
```

```
## [1] 4.0 5.0 2.0 5.0 3.0 15.0 3.0 2.0 0.0 2.0
## [15] 0.0 0.0 NA 3.0 0.0 4.0 3.0 4.0 10.0 2.5
## [29] 2.0 4.0 2.0 4.0 4.0 0.0 5.0 0.0 2.0
```

► As a last step, we should go through and figure out if any of

Rebuilding our data

```
survey <- transform(survey.messy, TVhours = tv.hours.clean)  
str(survey)
```

```
## 'data.frame':    37 obs. of  6 variables:  
## $ Program      : Factor w/ 3 levels "MISM","Other",..  
## $ PriorExp     : Factor w/ 3 levels "Extensive experi..  
## $ Rexperience  : Factor w/ 3 levels "Basic competence..  
## $ OperatingSystem: Factor w/ 2 levels "Mac OS X","Windo..  
## $ TVhours      : num  4 5 2 5 3 15 3 2 0 2 ...  
## $ Editor       : Factor w/ 5 levels "Google Docs",... ..
```

► **Success!**

A different approach

- ▶ We can also handle this problem by setting `stringsAsFactors = FALSE` when importing our data.

```
survey.messy <- read.csv("survey_messy.csv", header=TRUE, s  
str(survey.messy)
```

```
## 'data.frame':    37 obs. of  6 variables:  
## $ Program      : chr  "PPM" "PPM" "Other" "PPM" ...  
## $ PriorExp     : chr  "Never programmed before" "Some  
## $ Rexperience  : chr  "Never used" "Basic competence"  
## $ OperatingSystem: chr  "Windows" "Mac OS X" "Mac OS X"  
## $ TVhours      : chr  "4" "~5" "2" "5or so" ...  
## $ Editor       : chr  "Microsoft Word" "Microsoft Word"
```

- ▶ Now everything is a character instead of a factor

One-line cleanup

- ▶ Let's clean up the TVhours column and cast it to numeric all in one command

```
survey <- transform(survey.messy,  
                    TVhours = as.numeric(gsub("[^0-9.]", "",  
str(survey)
```

```
## 'data.frame':    37 obs. of  6 variables:  
## $ Program      : chr  "PPM" "PPM" "Other" "PPM" ...  
## $ PriorExp     : chr  "Never programmed before" "Some  
## $ Rexperience  : chr  "Never used" "Basic competence"  
## $ OperatingSystem: chr  "Windows" "Mac OS X" "Mac OS X"  
## $ TVhours      : num  4 5 2 5 3 15 3 2 0 2 ...  
## $ Editor       : chr  "Microsoft Word" "Microsoft Word"
```

What about all those other character variables?

```
table(survey[["Program"]])
```

```
##
```

```
##  MISM Other  PPM
```

```
##    13     7   17
```

```
table(as.factor(survey[["Program"]]))
```

```
##
```

```
##  MISM Other  PPM
```

```
##    13     7   17
```

- ▶ Having factors coded as characters may be OK for many parts of our analysis

To be safe, let's fix things

```
# Figure out which columns are coded as characters  
chr.indexes <- sapply(survey, FUN = is.character)  
chr.indexes
```

```
##          Program          PriorExp      Rexperience Operati  
##          TRUE          TRUE          TRUE  
##          TVhours          Editor  
##          FALSE          TRUE
```

```
# Re-code all of the character columns to factors  
survey[chr.indexes] <- lapply(survey[chr.indexes], FUN = as
```

Here's the outcome

```
str(survey)
```

```
## 'data.frame':    37 obs. of  6 variables:
## $ Program      : Factor w/ 3 levels "MISM","Other",...
## $ PriorExp     : Factor w/ 3 levels "Extensive experi
## $ Rexperience   : Factor w/ 3 levels "Basic competence
## $ OperatingSystem: Factor w/ 2 levels "Mac OS X","Wind
## $ TVhours       : num  4 5 2 5 3 15 3 2 0 2 ...
## $ Editor        : Factor w/ 5 levels "Google Docs",...
```

► **Success!**

Another common problem

- ▶ When data is entered manually, misspellings and case changes are very common
- ▶ E.g., a column showing life support mechanism may look like,

```
life.support <- as.factor(c("dialysis", "Ventilation", "Dialysis"))
summary(life.support)
```

```
## dialysis Dialysis dyalysis nnone n
##          3          1          1          1
## ventilation Ventilation
##          1          1
```

```
summary(life.support)
```

##	dialysis	Dialysis	dyalysis	nnone	n
##	3	1	1	1	

What are all these [l/s/t/]apply() functions?

- ▶ These are all efficient ways of applying a function to margins of an array or elements of a list
- ▶ Before we talk about the details of `apply()` and its relatives, we should first understand loops
- ▶ **loops** are ways of iterating over data
- ▶ The `apply()` functions can be thought of as good *alternatives* to loops

For loops: a pair of examples

```
for(i in 1:4) {  
  print(i)  
}
```

```
## [1] 1  
## [1] 2  
## [1] 3  
## [1] 4
```

```
phrase <- "Good Night, "  
for(word in c("and", "Good", "Luck")) {  
  phrase <- paste(phrase, word)  
  print(phrase)  
}
```

```
## [1] "Good Night,  and"  
## [1] "Good Night,  and Good"  
## [1] "Good Night,  and Good Luck"
```

For loops: syntax

A **for loop** executes a chunk of code for every value of an **index variable** in an **index set**

- ▶ The basic syntax takes the form

```
# for(index.variable in index.set) {  
#   code to be repeated at every value of index.variable  
# }
```

- ▶ The index set is often a vector of integers, but can be more general

Example

```
index.set <- list(name="Michael", weight=185, is.male=TRUE)
for(i in index.set) {
  print(c(i, typeof(i)))
}
```

```
## [1] "Michael"    "character"
## [1] "185"        "double"
## [1] "TRUE"       "logical"
```

Example: Calculate sum of each column

```
fake.data <- matrix(rnorm(500), ncol=5) # create fake 100 x 5 matrix
head(fake.data,2) # print first two rows
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]
## [1,] -1.6446803 -2.1532042  0.5611550 -0.6246019 -0.8565111
## [2,] -0.8199759 -0.5727423 -0.4925805  0.9939846  0.8331111
```

```
col.sums <- numeric(ncol(fake.data)) # variable to store results
for(i in 1:nrow(fake.data)) {
  col.sums <- col.sums + fake.data[i,] # add ith observation to column sums
}
col.sums
```

```
## [1] -1.4895621 -4.4912187 -0.8146502 -7.2401811 -5.8186701
```

```
colSums(fake.data) # A better approach (see also colMeans())
```

```
## [1] -1.4895621 -4.4912187 -0.8146502 -7.2401811 -5.8186701
```

while loops

- ▶ **while loops** repeat a chunk of code while the specified condition remains true

```
day <- 1
num.days <- 365
while(day <= num.days) {
  day <- day + 1
}
```

- ▶ We won't really be using while loops in this class
- ▶ Just be aware that they exist, and that they may become useful to you at some point in your analytics career

The various `apply()` functions

Command	Description
---------	-------------

<code>apply(X, MARGIN, FUN)</code>	Obtain a vector/array/list by applying FUN along the specified MARGIN of an array or matrix X
------------------------------------	---

<code>lapply(X, FUN)</code>	Obtain a list by applying FUN to the elements of a list X
-----------------------------	---

<code>sapply(X)</code>	Simplified
------------------------	------------

Example: apply()

```
colMeans(fake.data)
```

```
## [1] -0.014895621 -0.044912187 -0.008146502 -0.072401811
```

```
apply(fake.data, MARGIN=2, FUN=mean) # MARGIN = 1 for rows
```

```
## [1] -0.014895621 -0.044912187 -0.008146502 -0.072401811
```

```
# Function that calculates proportion of vector indexes that are positive
```

```
propPositive <- function(x) mean(x > 0)
```

```
apply(fake.data, MARGIN=2, FUN=propPositive)
```

```
## [1] 0.50 0.53 0.48 0.46 0.51
```


Example: lapply(), sapply()

```
lapply(survey, is.factor) # Returns a list
```

```
## $Program  
## [1] TRUE  
##  
## $PriorExp  
## [1] TRUE  
##  
## $Rexperience  
## [1] TRUE  
##  
## $OperatingSystem  
## [1] TRUE  
##  
## $TVhours  
## [1] FALSE  
##  
## $Editor
```

Example: `apply()`, `lapply()`, `sapply()`

```
apply(cars, 2, FUN=mean) # Data frames are arrays
```

```
## speed  dist  
## 15.40 42.98
```

```
lapply(cars, FUN=mean) # Data frames are also lists
```

```
## $speed  
## [1] 15.4  
##  
## $dist  
## [1] 42.98
```

```
sapply(cars, FUN=mean) # sapply() is just simplified lapply()
```

```
## speed  dist  
## 15.40 42.98
```

Example: `tapply()`

- ▶ Think of `tapply()` as a generalized form of the `table()` function

```
library(MASS)
```

```
##
```

```
## Attaching package: 'MASS'
```

```
## The following object is masked _by_ 'GlobalEnv':
```

```
##
```

```
##      survey
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
##      select
```

```
# Get a count table, data broken down by Origin and DriveTrain
```

```
table(Cars93$Origin, Cars93$DriveTrain)
```

Example: tapply()

- ▶ Let's get the average horsepower by car Origin and Type

```
tapply(Cars93[["Horsepower"]], INDEX = Cars93[c("Origin", "Type"],
```

```
##              Type
## Origin      Compact      Large  Midsize      Small      Sporty
##   USA      117.4286 179.4545 153.5000  89.42857 166.5000 1
## non-USA 141.5556      NA 189.4167  91.78571 151.6667 1
```

- ▶ What's that NA doing there?

```
any(Cars93$Origin == "non-USA" & Cars93$Type == "Large")
```

```
## [1] FALSE
```

- ▶ None of the non-USA manufacturers produced Large cars!

Example: using `tapply()` to mimic `table()`

- Here's how one can use `tapply()` to produce the same output as the `table()` function

```
library(MASS)
# Get a count table, data broken down by Origin and DriveTrain
table(Cars93$Origin, Cars93$DriveTrain)
```

```
##
##           4WD Front Rear
##   USA           5    34    9
## non-USA         5    33    7
```

```
# This one may take a moment to figure out...
tapply(rep(1, nrow(Cars93)), INDEX = Cars93[c("Origin", "DriveTrain")], FUN = function(x) length(x))
```

```
##           DriveTrain
## Origin    4WD Front Rear
##   USA           5    34    9
```

with()

- ▶ Thus far we've repeatedly typed out the data frame name when referencing its columns
- ▶ This is because the data variables don't exist in our working environment
- ▶ Using **with**(data, expr) lets us specify that the code in expr should be evaluated in an environment that contains the elements of data as variables

```
with(Cars93, table(Origin, Type))
```

```
##           Type
## Origin   Compact Large Midsize Small Sporty Van
##   USA           7    11      10     7      8    5
## non-USA        9     0      12    14      6    4
```

Example: with()

```
any(Cars93$Origin == "non-USA" & Cars93$Type == "Large")
```

```
## [1] FALSE
```

```
with(Cars93, any(Origin == "non-USA" & Type == "Large")) #
```

```
## [1] FALSE
```

```
with(Cars93, tapply(Horsepower, INDEX = list(Origin, Type),
```

```
##           Compact      Large  Midsize    Small    Sporty
## USA      117.4286 179.4545 153.5000  89.42857 166.5000 158
## non-USA 141.5556      NA 189.4167  91.78571 151.6667 138
```

- Using with() makes code simpler, easier to read, and easier to debug